

# ParamInsight: Bayesian Parameter Inference via MCMC

Kevin Mota da Costa

January 12, 2026

## Abstract

This document presents the theoretical background and implementation details of *ParamInsight*, a Python framework for Bayesian parameter inference using Markov Chain Monte Carlo (MCMC) methods. The project emphasizes statistical consistency, explicit likelihood modeling, realistic observational uncertainties, and a full MCMC implementation from scratch.

## 1 Introduction

Parameter inference aims to estimate model parameters given noisy observational data while properly accounting for uncertainty. *ParamInsight* provides a minimal yet rigorous framework that demonstrates:

- Likelihood-based Bayesian inference,
- Explicit construction of Gaussian noise using the Box–Muller transform,
- MCMC sampling via the Metropolis–Hastings algorithm,
- Realistic modeling of observational uncertainties,
- Statistical analysis and visualization of posterior distributions.

All components are implemented explicitly to highlight the underlying statistical and computational principles.

## 2 Observational Model and Likelihood

Let  $x_i$  denote the independent variable,  $y_i$  the observed data, and  $\delta y_i$  the uncertainty associated with each observation. For a generic two-parameter model  $F(x; a, b)$ , the log-likelihood function is defined as

$$\log L(a, b) = -\frac{1}{2} \sum_{i=1}^N \left( \frac{F(x_i; a, b) - y_i}{\delta y_i} \right)^2. \quad (1)$$

This corresponds to a chi-squared likelihood assuming independent Gaussian errors with *point-dependent* variances. The normalization constant of the Gaussian distribution is omitted since it does not affect likelihood ratios used by the MCMC acceptance rule.

### 3 Observational Uncertainty Model

Unlike fixed-error models, *ParamInsight* generates realistic observational uncertainties for each data point:

$$\delta y_i = |\delta_{\text{inst}} + c_{\text{trend}} x_i + \epsilon_i|, \quad (2)$$

where:

- $\delta_{\text{inst}}$  is a base instrumental error,
- $c_{\text{trend}}$  introduces a systematic trend with  $x$ ,
- $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$  is random observational noise.

Observed data are generated as

$$y_i = F(x_i; a_{\text{true}}, b_{\text{true}}) + \mathcal{N}(0, \delta y_i^2). \quad (3)$$

This approach mimics realistic measurement processes encountered in experimental and observational sciences.

## 4 Random Number Generation

### 4.1 Uniform Distribution

Uniform random numbers are generated using NumPy's core generator and mapped to arbitrary intervals.

### 4.2 Gaussian Distribution: Box–Muller Transformation

Gaussian random variables are generated explicitly using the Box–Muller transform:

$$z = \sqrt{-2 \ln u_1} \cos(2\pi u_2), \quad (4)$$

$$u_1, u_2 \sim \text{Uniform}(0, 1). \quad (5)$$

The resulting variable  $z$  follows a standard normal distribution and is scaled to arbitrary mean  $\mu$  and standard deviation  $\sigma$  when required. This implementation avoids reliance on black-box Gaussian generators.

## 5 MCMC Sampling Method

### 5.1 Metropolis–Hastings Algorithm

The parameter vector  $\theta = (a, b)$  is sampled using the Metropolis–Hastings algorithm. Given the current state  $\theta_{n-1}$ , a proposal  $\theta'$  is generated and accepted with probability:

$$\alpha = \min(1, \exp [\log L(\theta') - \log L(\theta_{n-1})]). \quad (6)$$

Rejected proposals result in the chain retaining its previous state.

## 5.2 Memory and Inertia Term

To improve exploration efficiency, *ParamInsight* introduces a simple inertia term using the two previous steps:

$$a' = a_{n-1} + \frac{1}{2}(a_{n-1} - a_{n-2}) + \eta_a, \quad (7)$$

$$b' = b_{n-1} + \frac{1}{2}(b_{n-1} - b_{n-2}) + \eta_b, \quad (8)$$

where  $\eta_a, \eta_b \sim \mathcal{N}(0, \sigma^2)$ . Although this proposal is not strictly Markovian, detailed balance is preserved through the acceptance rule, ensuring correct posterior sampling.

## 6 Implemented Models

The framework includes four example models:

- Linear:  $F(x) = ax + b$
- Logarithmic:  $F(x) = a \log(bx)$
- Quadratic:  $F(x) = ax + bx^2$
- Inverse:  $F(x) = a/x + b$

Each model is tested with synthetic observational data and independently sampled using the same MCMC engine.

## 7 Implementation Overview

The project is organized into the following modules:

- `distributions.py`: likelihood computation and random generators,
- `mcmc.py`: 2D Metropolis–Hastings sampler with inertia,
- `examples.py`: data generation, experiment setup, and execution,
- `utils.py`: visualization and result storage,
- `main.py`: entry point for running all experiments.

Each experiment saves observational data, MCMC chains, diagnostic plots, and final statistical summaries.

## 8 Conclusion

*ParamInsight* provides a transparent and statistically grounded framework for Bayesian parameter inference. By explicitly implementing likelihoods, noise models, and MCMC sampling from first principles, the project offers both educational value and a solid foundation for extension to more complex inference problems.

## 9 References

1. Gelman et al., *Bayesian Data Analysis*, CRC Press (2013).
2. Robert & Casella, *Monte Carlo Statistical Methods*, Springer (2004).
3. Box & Muller (1958), *Annals of Mathematical Statistics*.
4. Metropolis et al. (1953), *Journal of Chemical Physics*.
5. Neal (1993), Technical Report CRG-TR-93-1.